


Tea Leaf Classification



```
import numpy as np
import pandas as pd
import os

base_path = "/kaggle/input/tealeafbd-tea-leaf-disease-
detection/teaLeafBD/teaLeafBD"
categories = ["1. Tea algal leaf spot", "2. Brown Blight", "3. Gray Blight",
"4. Helopeltis" , "5. Red spider", "6. Green mirid bug", "7. Healthy leaf"]

image_paths = []
labels = []

for category in categories:
    category_path = os.path.join(base_path, category)
    for image_name in os.listdir(category_path):
        image_path = os.path.join(category_path, image_name)
        image_paths.append(image_path)
        labels.append(category)

df = pd.DataFrame({
    "image_path": image_paths,
    "label": labels
})

df.head()
```

```

                                image_path      label
0  /kaggle/input/tealeafbd-tea-leaf-disease-detec...  1. Tea algal leaf spot
1  /kaggle/input/tealeafbd-tea-leaf-disease-detec...  1. Tea algal leaf spot
2  /kaggle/input/tealeafbd-tea-leaf-disease-detec...  1. Tea algal leaf spot
3  /kaggle/input/tealeafbd-tea-leaf-disease-detec...  1. Tea algal leaf spot
4  /kaggle/input/tealeafbd-tea-leaf-disease-detec...  1. Tea algal leaf spot

```

```
df.tail()
```

```

                                image_path      label
5271 /kaggle/input/tealeafbd-tea-leaf-disease-detec...  7. Healthy leaf
5272 /kaggle/input/tealeafbd-tea-leaf-disease-detec...  7. Healthy leaf
5273 /kaggle/input/tealeafbd-tea-leaf-disease-detec...  7. Healthy leaf
5274 /kaggle/input/tealeafbd-tea-leaf-disease-detec...  7. Healthy leaf
5275 /kaggle/input/tealeafbd-tea-leaf-disease-detec...  7. Healthy leaf

```

```
df.shape
```

```
(5276, 2)
```

```
df.columns
```

```
Index(['image_path', 'label'], dtype='object')
```

```
df.duplicated().sum()
```

```
0
```

```
df.isnull().sum()
```

```

image_path    0
label         0
dtype: int64

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5276 entries, 0 to 5275
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0  image_path  5276 non-null   object
 1  label       5276 non-null   object
dtypes: object(2)
memory usage: 82.6+ KB

```

```
df['label'].unique()
```

```

array(['1. Tea algal leaf spot', '2. Brown Blight', '3. Gray Blight',
      '4. Helopeltis', '5. Red spider', '6. Green mirid bug',
      '7. Healthy leaf'], dtype=object)

```

```
df['label'].value_counts()
```

```
label
6. Green mirid bug      1282
3. Gray Blight         1013
7. Healthy leaf        935
4. Helopeltis          607
5. Red spider          515
2. Brown Blight        506
1. Tea algal leaf spot 418
Name: count, dtype: int64
```

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
sns.set_style("whitegrid")
```

```
fig, ax = plt.subplots(figsize=(8, 6))
sns.countplot(data=df, x="label", palette="viridis", ax=ax)
```

```
ax.set_title("Distribution of Disease Types", fontsize=14, fontweight='bold')
ax.set_xlabel("Tumor Type", fontsize=12)
ax.set_ylabel("Count", fontsize=12)
```

```
for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}',
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='bottom', fontsize=11, color='black',
                xytext=(0, 5), textcoords='offset points')
```

```
plt.xticks(rotation=-45)
plt.show()
```

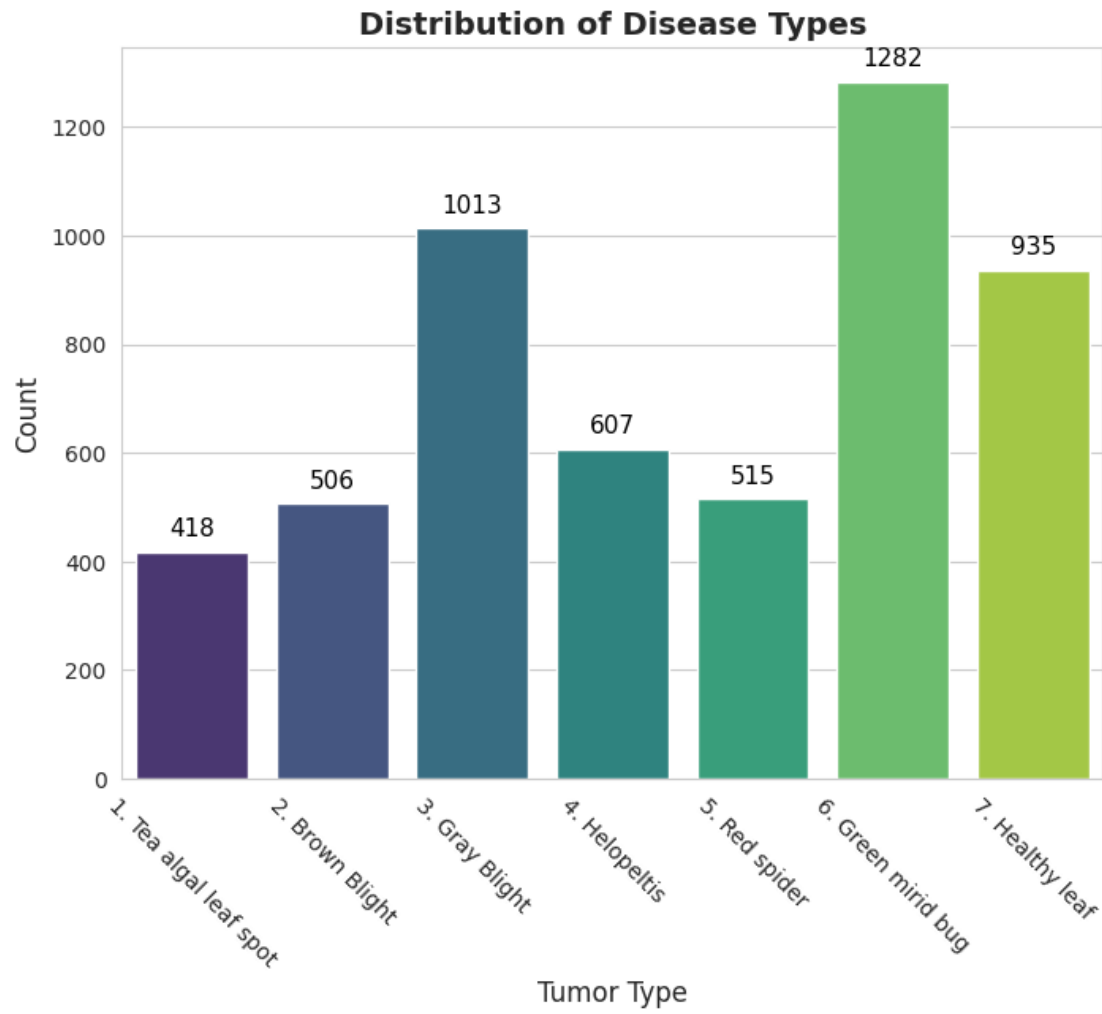
```
label_counts = df["label"].value_counts()
```

```
fig, ax = plt.subplots(figsize=(8, 6))
colors = sns.color_palette("viridis", len(label_counts))
```

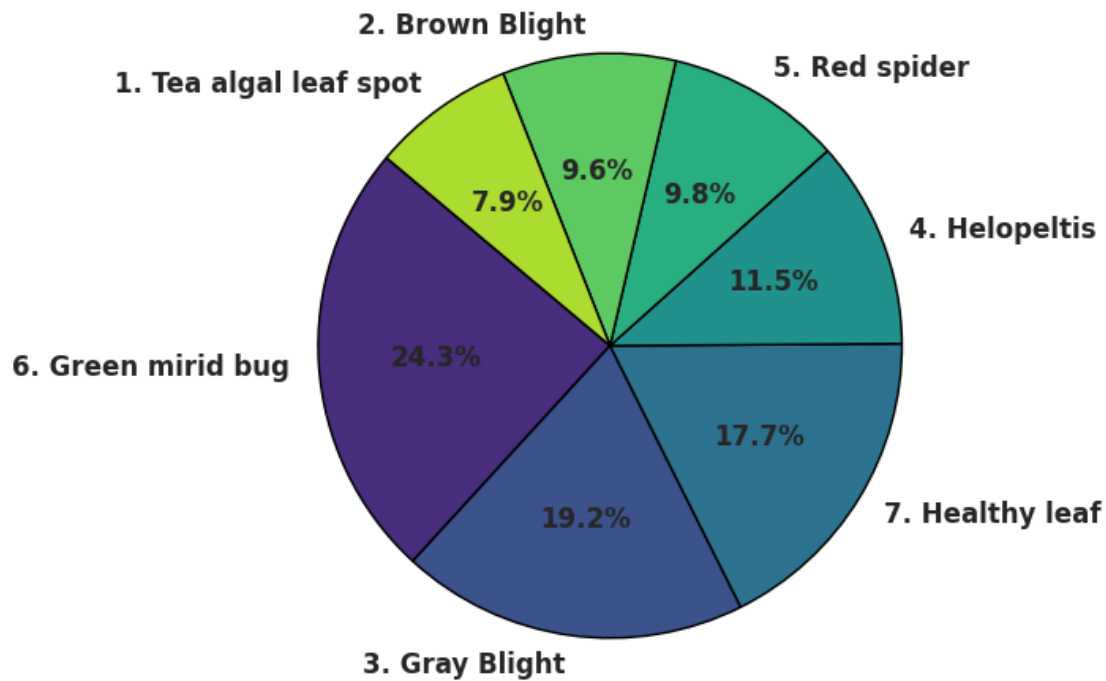
```
ax.pie(label_counts, labels=label_counts.index, autopct='%1.1f%%',
        startangle=140, colors=colors, textprops={'fontsize': 12, 'weight':
'bold'},
        wedgeprops={'edgecolor': 'black', 'linewidth': 1})
```

```
ax.set_title("Distribution of Disease Types - Pie Chart", fontsize=14,
fontweight='bold')
```

```
plt.show()
```



Distribution of Disease Types - Pie Chart



```
import cv2

num_images = 5

plt.figure(figsize=(15, 12))

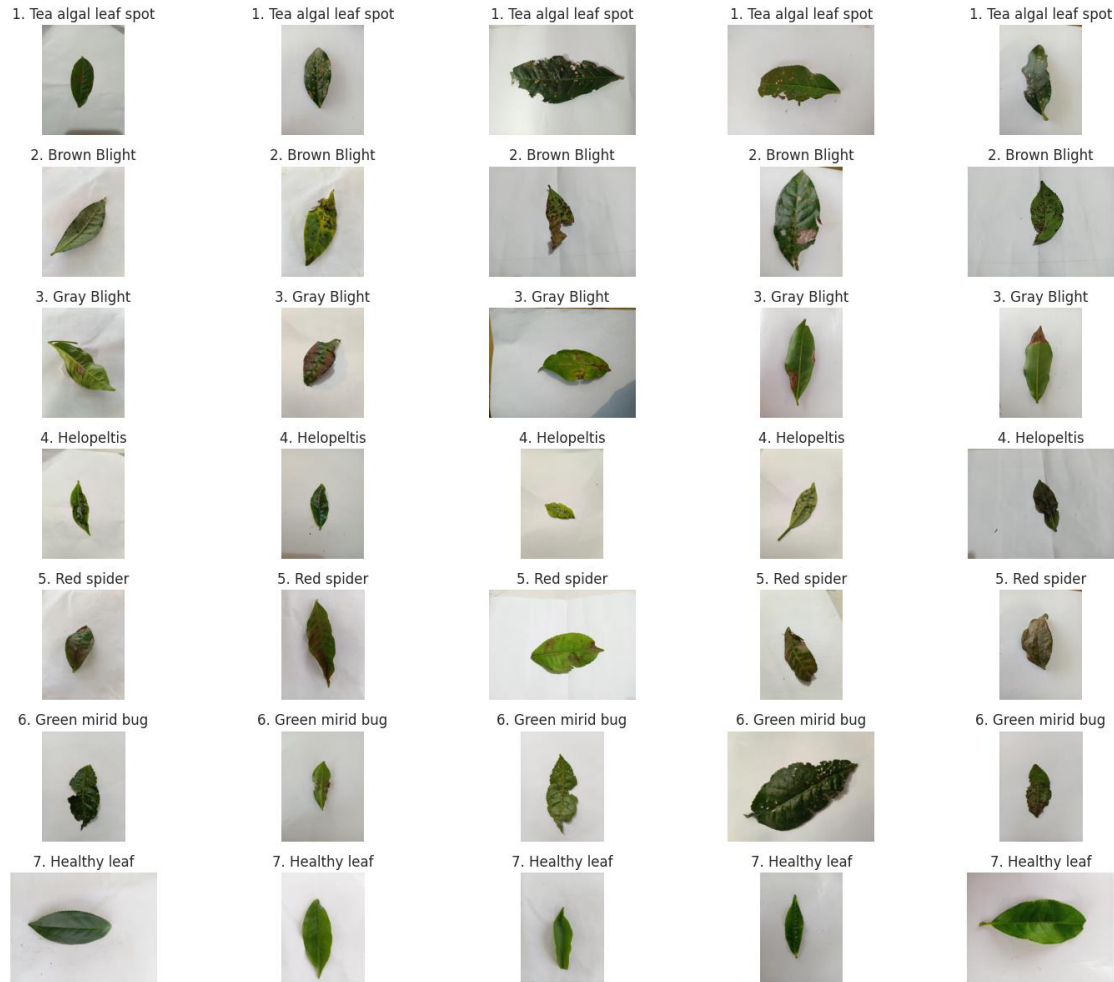
for i, category in enumerate(categories):
    category_images = df[df['label'] ==
category]['image_path'].iloc[:num_images]

    for j, img_path in enumerate(category_images):

        img = cv2.imread(img_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        plt.subplot(len(categories), num_images, i * num_images + j + 1)
        plt.imshow(img)
        plt.axis('off')
        plt.title(category)

plt.tight_layout()
plt.show()
```



```

from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
df['category_encoded'] = label_encoder.fit_transform(df['label'])

df = df[['image_path', 'category_encoded']]

from sklearn.utils import resample

max_count = df['category_encoded'].value_counts().max()

dfs = []
for category in df['category_encoded'].unique():
    class_subset = df[df['category_encoded'] == category]
    class_upsampled = resample(class_subset,
                              replace=True,
                              n_samples=max_count,
                              random_state=42)
    dfs.append(class_upsampled)

```

```
df_balanced = pd.concat(dfs).sample(frac=1,
random_state=42).reset_index(drop=True)
```

```
df_balanced['category_encoded'].value_counts()
```

```
category_encoded
```

```
3    1282
```

```
0    1282
```

```
1    1282
```

```
2    1282
```

```
4    1282
```

```
5    1282
```

```
6    1282
```

```
Name: count, dtype: int64
```

```
df_resampled = df_balanced
```

```
df_resampled['category_encoded'] =
df_resampled['category_encoded'].astype(str)
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
import tensorflow as tf
```

```
from tensorflow import keras
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.optimizers import Adam
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
Activation, Dropout, BatchNormalization
```

```
from tensorflow.keras import regularizers
```

```
import warnings
```

```
warnings.filterwarnings("ignore")
```

```
print ('check')
```

```
2025-06-09 06:31:21.834484: E
```

```
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to
register cuFFT factory: Attempting to register factory for plugin cuFFT when
one has already been registered
```

```
WARNING: All log messages before absl::InitializeLog() is called are written
to STDERR
```

```
E0000 00:00:1749450682.331013      35 cuda_dnn.cc:8310] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one has
already been registered
```

```
E0000 00:00:1749450682.459893      35 cuda_blas.cc:1418] Unable to register
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has
already been registered
```

```
check
```

```

train_df_new, temp_df_new = train_test_split(
    df_resampled,
    train_size=0.8,
    shuffle=True,
    random_state=42,
    stratify=df_resampled['category_encoded']
)

valid_df_new, test_df_new = train_test_split(
    temp_df_new,
    test_size=0.5,
    shuffle=True,
    random_state=42,
    stratify=temp_df_new['category_encoded']
)

from tensorflow.keras.preprocessing.image import ImageDataGenerator

batch_size = 16
img_size = (224, 224)
channels = 3
img_shape = (img_size[0], img_size[1], channels)

tr_gen = ImageDataGenerator(
    rescale=1./255
)

ts_gen = ImageDataGenerator(rescale=1./255)

train_gen_new = tr_gen.flow_from_dataframe(
    train_df_new,
    x_col='image_path',
    y_col='category_encoded',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)

valid_gen_new = ts_gen.flow_from_dataframe(
    valid_df_new,
    x_col='image_path',
    y_col='category_encoded',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)

```



```
)

test_gen_new = ts_gen.flow_from_dataframe(
    test_df_new,
    x_col='image_path',
    y_col='category_encoded',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=False,
    batch_size=batch_size
)
```

Found 7179 validated image filenames belonging to 7 classes.

Found 897 validated image filenames belonging to 7 classes.

Found 898 validated image filenames belonging to 7 classes.

```
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

Num GPUs Available: 2

```
gpus = tf.config.list_physical_devices('GPU')
if gpus:
    try:
        for gpu in gpus:
            tf.config.experimental.set_memory_growth(gpu, True)
        print("GPU is set for TensorFlow")
    except RuntimeError as e:
        print(e)
```

GPU is set for TensorFlow

```
from tensorflow.keras import layers, models
```

```
num_classes = 7
```

```
class ContinuousLayer(layers.Layer):
    def __init__(self, kernel_size=5, num_basis=10, output_channels=16,
**kwargs):
        super(ContinuousLayer, self).__init__(**kwargs)
        self.kernel_size = kernel_size
        self.num_basis = num_basis
        self.output_channels = output_channels
        self.centers = self.add_weight(
            name='centers',
            shape=(num_basis, 2),
            initializer='random_normal',
            trainable=True
        )
        self.widths = self.add_weight(
            name='widths',
            shape=(num_basis,),
```

```

        initializer='ones',
        trainable=True,
        constraint=tf.keras.constraints.NonNeg()
    )
    self.kernel_weights = self.add_weight(
        name='kernel_weights',
        shape=(kernel_size, kernel_size, channels, output_channels),
        initializer='glorot_normal',
        trainable=True
    )

    def call(self, inputs):
        height, width = img_size
        x = tf.range(0, height, 1.0)
        y = tf.range(0, width, 1.0)
        x_grid, y_grid = tf.meshgrid(x, y)
        grid = tf.stack([x_grid, y_grid], axis=-1)

        basis = []
        for i in range(self.num_basis):
            center = self.centers[i]
            width = self.widths[i]
            dist = tf.reduce_sum(((grid - center) / width) ** 2, axis=-1)
            basis_i = tf.exp(-dist)
            basis.append(basis_i)
        basis = tf.stack(basis, axis=-1)

        basis_weights = tf.reduce_mean(basis, axis=[0, 1])
        basis_weights = tf.nn.softmax(basis_weights)
        basis_weights = basis_weights[:, tf.newaxis, tf.newaxis, tf.newaxis,
tf.newaxis]

        modulated_kernel = self.kernel_weights * tf.reduce_sum(basis_weights,
axis=0)

        output = tf.nn.conv2d(
            inputs,
            modulated_kernel,
            strides=[1, 1, 1, 1],
            padding='SAME'
        )

        return output

    def compute_output_shape(self, input_shape):
        return (input_shape[0], input_shape[1], input_shape[2],
self.output_channels)

    def smoothness_penalty(self):

```

```

        grad_x = tf.reduce_mean(tf.square(self.kernel_weights[1:, :, :, :] -
self.kernel_weights[:-1, :, :, :]))
        grad_y = tf.reduce_mean(tf.square(self.kernel_weights[:, 1:, :, :] -
self.kernel_weights[:, :-1, :, :]))
        return grad_x + grad_y

class VariationalLoss(tf.keras.losses.Loss):
    def __init__(self, model, lambda1=0.01, lambda2=1.0):
        super(VariationalLoss, self).__init__()
        self.model = model
        self.lambda1 = lambda1
        self.lambda2 = lambda2
        self.sce = tf.keras.losses.SparseCategoricalCrossentropy() # Changed
to SparseCategoricalCrossentropy

    def call(self, y_true, y_pred):
        smoothness_penalty = 0
        for layer in self.model.layers:
            if isinstance(layer, ContinuousLayer):
                smoothness_penalty += layer.smoothness_penalty()
        prediction_loss = self.sce(y_true, y_pred)
        return self.lambda2 * prediction_loss + self.lambda1 *
smoothness_penalty

def build_continuous_model():
    inputs = layers.Input(shape=img_shape)
    x = ContinuousLayer(kernel_size=5, num_basis=10,
output_channels=16)(inputs)
    x = layers.Activation('relu')(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.Flatten()(x)
    x = layers.Dense(128, activation='relu')(x)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x) # Changed
to 7 units with softmax
    model = models.Model(inputs, outputs)
    return model

model = build_continuous_model()

model.compile(
    optimizer='adam',
    loss=VariationalLoss(model=model, lambda1=0.01, lambda2=1.0),
    metrics=['accuracy']
)

history = model.fit(
    train_gen_new,
    validation_data=valid_gen_new,

```

```
epochs=3,  
verbose=1  
)
```

Epoch 1/3

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

```
I0000 00:00:1749451063.329868    141 service.cc:148] XLA service  
0x7da2d8025740 initialized for platform CUDA (this does not guarantee that  
XLA will be used). Devices:  
I0000 00:00:1749451063.331201    141 service.cc:156]   StreamExecutor device  
(0): Tesla T4, Compute Capability 7.5  
I0000 00:00:1749451063.331220    141 service.cc:156]   StreamExecutor device  
(1): Tesla T4, Compute Capability 7.5  
I0000 00:00:1749451063.860056    141 cuda_dnn.cc:529] Loaded cuDNN version  
90300
```

```
1/449 _____ 1:15:43 10s/step - accuracy: 0.1875 - loss:  
1.9956
```

```
I0000 00:00:1749451069.707228    141 device_compiler.h:188] Compiled cluster  
using XLA! This line is logged at most once for the lifetime of the process.
```

```
449/449 _____ 138s 285ms/step - accuracy: 0.2923 - loss:  
2.3807 - val_accuracy: 0.5953 - val_loss: 1.1260
```

Epoch 2/3

```
449/449 _____ 83s 185ms/step - accuracy: 0.6487 - loss: 1.0014  
- val_accuracy: 0.7915 - val_loss: 0.6917
```

Epoch 3/3

```
449/449 _____ 84s 187ms/step - accuracy: 0.8393 - loss: 0.5007  
- val_accuracy: 0.8506 - val_loss: 0.4932
```

model.summary()

Model: "functional"

Layer (type) Param #	Output Shape	
input_layer_2 (InputLayer) 0	(None, 224, 224, 3)	
continuous_layer_2 (ContinuousLayer) 1,230	(None, 224, 224, 16)	
activation_2 (Activation)	(None, 224, 224, 16)	

0				
		max_pooling2d_2 (MaxPooling2D)	(None, 112, 112, 16)	
0				
		flatten_2 (Flatten)	(None, 200704)	
0				
		dense_3 (Dense)	(None, 128)	
25,690,240				
		dropout_2 (Dropout)	(None, 128)	
0				
		dense_4 (Dense)	(None, 7)	
903				

Total params: 77,077,121 (294.03 MB)

Trainable params: 25,692,373 (98.01 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 51,384,748 (196.02 MB)

```
test_loss, test_accuracy = model.evaluate(test_gen_new)
print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy:.4f}")
```



57/57 ————— 12s 212ms/step - accuracy: 0.8856 - loss: 0.4414
Test Loss: 0.4905, Test Accuracy: 0.8719

```
import matplotlib.pyplot as plt

plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
```

```
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```




```
test_gen_new.reset()
y_pred = model.predict(test_gen_new)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = test_gen_new.classes
```

57/57 ————— 9s 158ms/step

```
cm = confusion_matrix(y_true, y_pred_classes)
```

```
class_names = list(test_gen_new.class_indices.keys())
```

```
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names,
yticklabels=class_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```



```

class ContinuousLayer(layers.Layer):
    def __init__(self, kernel_size=5, num_basis=10, output_channels=16,
reduction_ratio=4, **kwargs):
        super(ContinuousLayer, self).__init__(**kwargs)
        self.kernel_size = kernel_size
        self.num_basis = num_basis
        self.output_channels = output_channels
        self.reduction_ratio = reduction_ratio

        self.centers = self.add_weight(
            name='centers',
            shape=(num_basis, 2),
            initializer='random_normal',
            trainable=True
        )
        self.widths = self.add_weight(
            name='widths',
            shape=(num_basis,),
            initializer='ones',
            trainable=True,

```

```

        constraint=tf.keras.constraints.NonNeg()
    )

    self.kernel_weights = self.add_weight(
        name='kernel_weights',
        shape=(kernel_size, kernel_size, channels, output_channels),
        initializer='glorot_normal',
        trainable=True
    )

    self.attention_pooling = layers.GlobalAveragePooling2D()
    self.attention_dense1 = layers.Dense(
        max(num_basis // self.reduction_ratio, 1),
        activation='relu',
        name='attention_dense1'
    )

    self.attention_dense2 = layers.Dense(
        num_basis,
        activation='sigmoid',
        name='attention_dense2'
    )

def call(self, inputs):
    height, width = img_size

    x_coords = tf.range(0, height, 1.0)
    y_coords = tf.range(0, width, 1.0)
    x_grid, y_grid = tf.meshgrid(x_coords, y_coords)
    grid = tf.stack([x_grid, y_grid], axis=-1)

    centers_resaped = self.centers[tf.newaxis, tf.newaxis, :, :]
    widths_resaped = self.widths[tf.newaxis, tf.newaxis, :, tf.newaxis]
    safe_widths = tf.maximum(widths_resaped, tf.keras.backend.epsilon())

    diff = (grid[:, :, tf.newaxis, :] - centers_resaped) / safe_widths
    dist_squared = tf.reduce_sum(diff ** 2, axis=-1)
    basis = tf.exp(-dist_squared)

    squeeze = self.attention_pooling(inputs)

    excitation = self.attention_dense1(squeeze)
    attention_weights = self.attention_dense2(excitation)

    mean_basis_activation = tf.reduce_mean(basis, axis=[0, 1]) #
(num_basis,)

    dynamic_basis_modulation = attention_weights *
    mean_basis_activation[tf.newaxis, :] # (batch_size, num_basis)

```



```

        scaling_factor_per_batch = tf.reduce_sum(dynamic_basis_modulation,
axis=-1, keepdims=True) # (batch_size, 1)
        global_features = self.attention_pooling(inputs) # (batch_size,
input_channels)

        predicted_basis_weights_per_batch =
self.attention_dense1(global_features)
        predicted_basis_weights_per_batch =
self.attention_dense2(predicted_basis_weights_per_batch) # (batch_size,
num_basis)

        attended_basis_weights =
tf.reduce_mean(predicted_basis_weights_per_batch, axis=0) # (num_basis,)

        attended_basis_weights = tf.nn.softmax(attended_basis_weights)

        scaling_factor = tf.reduce_sum(attended_basis_weights)

        modulated_kernel = self.kernel_weights * scaling_factor

        output = tf.nn.conv2d(
            inputs,
            modulated_kernel,
            strides=[1, 1, 1, 1],
            padding='SAME'
        )

        return output

    def compute_output_shape(self, input_shape):
        return (input_shape[0], input_shape[1], input_shape[2],
self.output_channels)

    def smoothness_penalty(self):
        grad_x = tf.reduce_mean(tf.square(self.kernel_weights[1:, :, :, :] -
self.kernel_weights[:-1, :, :, :]))
        grad_y = tf.reduce_mean(tf.square(self.kernel_weights[:, 1:, :, :] -
self.kernel_weights[:, :-1, :, :]))
        return grad_x + grad_y

class VariationalLoss(tf.keras.losses.Loss):
    def __init__(self, model, lambda1=0.01, lambda2=1.0):
        super(VariationalLoss, self).__init__()
        self.model = model
        self.lambda1 = lambda1
        self.lambda2 = lambda2
        self.sce = tf.keras.losses.SparseCategoricalCrossentropy()

```

```

def call(self, y_true, y_pred):
    smoothness_penalty = 0
    for layer in self.model.layers:
        if isinstance(layer, ContinuousLayer):
            smoothness_penalty += layer.smoothness_penalty()
    prediction_loss = self.sce(y_true, y_pred)
    return self.lambda2 * prediction_loss + self.lambda1 *
smoothness_penalty

def build_continuous_model_with_attention():
    inputs = layers.Input(shape=img_shape)
    x = ContinuousLayer(kernel_size=5, num_basis=10, output_channels=16,
reduction_ratio=4)(inputs)
    x = layers.Activation('relu')(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.Flatten()(x)
    x = layers.Dense(128, activation='relu')(x)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)
    model = models.Model(inputs, outputs)
    return model

```

```
model.summary()
```

```
Model: "functional_4"
```

Layer (type) Param #	Output Shape	
input_layer_6 (InputLayer) 0	(None, 224, 224, 3)	
continuous_layer_8 (ContinuousLayer) 1,268	(None, 224, 224, 16)	
activation_8 (Activation) 0	(None, 224, 224, 16)	
max_pooling2d_8 (MaxPooling2D) 0	(None, 112, 112, 16)	
flatten_6 (Flatten) 0	(None, 200704)	

dense_11 (Dense)	(None, 128)	
25,690,240		
dropout_6 (Dropout)	(None, 128)	
0		
dense_12 (Dense)	(None, 7)	
903		

Total params: 77,077,235 (294.03 MB)

Trainable params: 25,692,411 (98.01 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 51,384,824 (196.02 MB)

```
model = build_continuous_model()
```

```
model.compile(
    optimizer='adam',
    loss=VariationalLoss(model=model, lambda1=0.01, lambda2=1.0),
    metrics=['accuracy']
)
```

```
history = model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=10,
    verbose=1
)
```

Epoch 1/10

449/449 ————— 91s 193ms/step - accuracy: 0.2086 - loss: 5.5501
- val_accuracy: 0.5530 - val_loss: 1.4257

Epoch 2/10

449/449 ————— 83s 185ms/step - accuracy: 0.5187 - loss: 1.3297
- val_accuracy: 0.7402 - val_loss: 0.8669

Epoch 3/10

449/449 ————— 83s 186ms/step - accuracy: 0.7466 - loss: 0.7435
- val_accuracy: 0.8384 - val_loss: 0.5552

Epoch 4/10

449/449 ————— 84s 187ms/step - accuracy: 0.8535 - loss: 0.4453
- val_accuracy: 0.8629 - val_loss: 0.4652

Epoch 5/10

```

449/449 ————— 83s 186ms/step - accuracy: 0.9069 - loss: 0.2940
- val_accuracy: 0.8640 - val_loss: 0.4886
Epoch 6/10
449/449 ————— 80s 179ms/step - accuracy: 0.9330 - loss: 0.2144
- val_accuracy: 0.8740 - val_loss: 0.4286
Epoch 7/10
449/449 ————— 84s 187ms/step - accuracy: 0.9487 - loss: 0.1594
- val_accuracy: 0.8829 - val_loss: 0.4688
Epoch 8/10
449/449 ————— 81s 181ms/step - accuracy: 0.9430 - loss: 0.1879
- val_accuracy: 0.8807 - val_loss: 0.4741
Epoch 9/10
449/449 ————— 83s 186ms/step - accuracy: 0.9462 - loss: 0.1729
- val_accuracy: 0.8629 - val_loss: 0.5662
Epoch 10/10
449/449 ————— 84s 187ms/step - accuracy: 0.9525 - loss: 0.1502
- val_accuracy: 0.8763 - val_loss: 0.5499

```

```



plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

```

```

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

```



```

test_gen_new.reset()
y_pred = model.predict(test_gen_new)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = test_gen_new.classes

57/57 ————— 10s 166ms/step

cm = confusion_matrix(y_true, y_pred_classes)

class_names = list(test_gen_new.class_indices.keys())

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names,
yticklabels=class_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

```

